

Using topic models as classifiers

TOPIC MODELING IN R



Pavel Oleinikov

Associate Director, Quantitative Analysis
Center, Wesleyan University

Topic models as soft classifiers

- Classification task - find probability of belonging to class
 - Named entity recognition - is an entity a geographic name or a person?
 - "***Washington*** crossed the Delaware river" vs. "They did a road trip across ***Washington***"
- Topic modeling of context - a set of words occurring next to the entity

Effect of control parameter alpha

- $k = 2$, $\alpha = 1$

	document	`1`	`2`
	<chr>	<dbl>	<dbl>
1	d_1	0.154	0.846
2	d_2	0.278	0.722
3	d_3	0.875	0.125
4	d_4	0.923	0.0769
5	d_5	0.5	0.5

- $k = 2$, $\alpha = \frac{50}{k} = 25$

	document	`1`	`2`
	<chr>	<dbl>	<dbl>
1	d_1	0.475	0.525
2	d_2	0.530	0.470
3	d_3	0.482	0.518
4	d_4	0.508	0.492
5	d_5	0.5	0.5

How LDA fits a model


- A bag of M&M candy - multinomial distribution
 - Several outcomes (colors of candy) repeated n times
 - Each outcome has its own probability - fixed by the factory that filled the bag
 - Probabilities sum up to 1
 - What's the probability of getting 5 yellow, 2 brown, 2 blue, and 1 black when we take out 10 pieces of candy?
- In LDA model, topics are color, and there are two "bags of candy": one for documents and one for words

The Dirichlet in LDA

- Randomized search for the best values of probabilities
 - E.g., try 0.2, 0.5 and 0.3 for proportions of three topics
- Hard to do for large number of topics and words
- Instead, use values from Dirichlet distribution as guesses
- Dirichlet distribution returns a set of numbers that add up to 1 - serve as probabilities of colors for M&M candy bags

```
      [,1] [,2] [,3]
[1,] 0.604 0.100 0.295
[2,] 0.133 0.609 0.259
[3,] 0.514 0.221 0.265
[4,] 0.113 0.112 0.775
[5,] 0.258 0.502 0.240
```



Dirichlet distribution



- Example: density profile of a 3-dimensional symmetric Dirichlet distribution
- Corners correspond to $(1,0,0)$, $(0,1,0)$, and $(0,0,1)$ combinations
- Favors assignment of a document to a single topic

alpha and the shape of Dirichlet distribution

- Left: $\alpha > 1$, right: $\alpha < 1$



Let's practice!
TOPIC MODELING IN R

From word windows to dtm

TOPIC MODELING IN R



Pavel Oleinikov

Associate Director, Quantitative Analysis
Center, Wesleyan University

Word window

- Entity is a personal noun (capitalized word)
- If we take n words on the left and n words on the right of an entity, we get a word window
- E.g., *attention of **Megara** was turned*
 - entity is Megara
 - context is *attention of was turned*
- Also possible to tag words to differentiate the side
 - attention_L1 of_L2 was_R1 turned_R2

A document for every entity

- Combine (using `paste`) contexts of the same entity to make a document

```
docs <- df %>% group_by(entity) %>%  
  summarise(doc_id = first(entity),  
            text = paste(text, collapse=" "))
```

- Example
 - entity - Anastasius
 - document (made of 2 word windows):
treasure_L1 which_L2 had_R1 bequeathed_R2 two_L1 years_L2 was_R1 overthrown_R2

Finding entities with regular expressions & stringr

- An entity is a capitalized word
- Find it using regular expressions pattern matching

```
pattern <- "[A-Z][a-z]+"  
m <- gregexpr(text, pattern)  
entities <- unlist(regmatches(text, m))
```

- Regular expressions:
 - character class `[A-Z]` and `[a-z]`
 - quantifiers: `{n,m}` - character occurs at least `n` times and at most `m` times
 - Quantifier shortcuts: `?` is `{0,1}`, `*` is `{0,}`, `+` is `{1,}`
 - `[A-Z][a-z]+` - one uppercase letter followed by one or more lowercase

Regular expressions with groups

- Parentheses serve to group some patterns together
 - Possible to "capture" the groups
- Some entities include `St.` in them, e.g. `St. Sophia`
- New pattern:

```
p <- "(St[.] )?[A-Z][a-z]+"
```

- `(St[.])` is a group. The `?` quantifier means the group is optional

Using capture groups to add a suffix

- Contents of a group can be back-referenced in substitution operations

```
t <- "the great Darius threw across"  
gsub("^[a-z]+) ([a-z]+)", "\\1_L1 \\2_L2", t)
```

- Two groups, each matches a lowercase word `[a-z]+`
- The `^` is an anchor - specifies position in the string. `^` - the start, `$` - at the end
- The `\\1` is back-reference to contents of group 1. Its contents are substituted.
- Result

```
"the_L1 great_L2 Darius threw across"
```

Let's practice

TOPIC MODELING IN R

Corpus alignment and classification

TOPIC MODELING IN R



Pavel Oleinikov

Associate Director, Quantitative Analysis
Center, Wesleyan University

Unsupervised classification

- Topic model for `k=3` useful for telling the meaning of a named entity

```
topics = tidy(mod, matrix="gamma") %>%  
  spread(topic, gamma)  
topics %>%  
  filter(document %in% c(" Alboin", " Alexander", " Asia Minor",  
    " Amorium", " Cappadocian"))
```

	document	`1`	`2`	`3`
	<chr>	<dbl>	<dbl>	<dbl>
1	" Alboin"	0.143	0.143	0.714
2	" Alexander"	0.143	0.143	0.714
3	" Amorium"	0.364	0.364	0.273
4	" Asia Minor"	0.0213	0.723	0.255
5	" Cappadocian"	0.571	0.143	0.286

Using pre-trained model

- Apply a pre-trained topic model to new data
- Function `posterior` in package `topicmodels`

```
model = LDA(...)  
result = posterior(model, new_data)  
result$topics
```

- `new_data` must be aligned with the vocabulary used in the model
- LDA algorithm iterates over items and their counts, does not "know" that it's dealing with words

Corpus alignment

- Drop words from dtm that are not part of model's vocabulary
- Function `tidy` with `matrix="beta"` extracts the terms and their probabilities

```
model_vocab <- tidy(mod, matrix="beta") %>%  
  select(term) %>% distinct()
```

- Do right-join with the model's vocabulary to keep only the words the model was trained on

```
new_table <- new_doc %>%  
  unnest_tokens(input=text, output=word) %>%  
  count(doc_id, word) %>%  
  right_join(model_vocab, by=c("word"="term"))
```

- Side effect - NA values

Handling NA values

- Right join assigns NA values in columns of rows for which there was no match

```
doc_id word      n
<chr>  <chr>    <int>
1 NA     emerged_r1 NA
2 NA     from_r2    NA
3 NA     horde_l1    NA
```

- We will end up with a new "document" - its name will be "NA"
- Solution: set NA counts to 0, set NA doc_id to the first "good" doc id

```
new_dtm <- new_table %>%
  arrange(desc(doc_id)) %>%
  mutate(doc_id = ifelse(is.na(doc_id), first(doc_id), doc_id),
         n = ifelse(is.na(n), 0, n)) %>%
  cast_dtm(document=doc_id, term=word, value=n)
```

Held-out data

- Machine learning: training vs. test data
- Held-out data for testing
 - Hold out a percentage of full records (same as with test datasets in ML)
 - Hold out a percentage of terms inside a document (unique to topic modeling)
- Estimate quality of fit by looking at the log-likelihood
 - "held-out log-likelihood"
- Our case: withhold full documents, no cross-validation

Let's practice!
TOPIC MODELING IN R